



Qwythos-gguf

 **Hugging Face**

Models | Datasets | Spaces | Buckets **new** | Docs | Enterprise | Pricing | 

Buckets:  Matuxy / **Qwythos-9B-Claude-Mythos-5-1M-GGUF-bucket** 

Files   Settings

0 Bytes | 15 files | Updated less than a minute ago



Name	Size	Uploaded	Xet hash
 .gitattributes	2.47 kB  	less than a minute ago	3f51bd66
 Qwythos-9B-Claude-Mythos-5-1M-BF16.gguf	17.9 GB  	less than a minute ago	4552ce52
 Qwythos-9B-Claude-Mythos-5-1M-MTP-BF16.gguf	18.4 GB  	less than a minute ago	a68d8b88
 Qwythos-9B-Claude-Mythos-5-1M-MTP-Q4_K_M.gguf	5.89 GB  	less than a minute ago	89294873
 Qwythos-9B-Claude-Mythos-5-1M-MTP-Q5_K_M.gguf	6.73 GB  	less than a minute ago	18f9ba76
 Qwythos-9B-Claude-Mythos-5-1M-MTP-Q6_K_M.gguf	7.62 GB  	less than a minute ago	599f762b
 Qwythos-9B-Claude-Mythos-5-1M-MTP-Q8_0.gguf	9.79 GB  	less than a minute ago	e8866ace
 Qwythos-9B-Claude-Mythos-5-1M-Q4_K_M.gguf	5.63 GB  	less than a minute ago	327ba44d
 Qwythos-9B-Claude-Mythos-5-1M-Q5_K_M.gguf	6.47 GB  	less than a minute ago	52e687cd
 Qwythos-9B-Claude-Mythos-5-1M-Q6_K_M.gguf	7.36 GB  	less than a minute ago	d447a5e1
 Qwythos-9B-Claude-Mythos-5-1M-Q8_0.gguf	9.53 GB  	less than a minute ago	3dffc996
 README.md	10.3 kB  	less than a minute ago	2982ae58
 SHA256SUMS	1.31 kB  	less than a minute ago	31ae6d93
 TEST_REPORT.md	3.72 kB  	less than a minute ago	2de3f342
 mmproj-Qwythos-9B-Claude-Mythos-5-1M-F16.gguf	918 MB  	less than a minute ago	58f3899f



Total size 0 Bytes


Files 15

Last updated Jul 25

Pre-warmed CDN    


Contributors





Qwythos-9B-Claude-Mythos-5-1M-GGUF

Developed by Matuxyz



GGUF quantizations of [empero-ai/Qwythos-9B-Claude-Mythos-5-1M](#) for [llama.cpp](#), Ollama, LM Studio, jan, KoboldCpp, and other GGUF runtimes.

Qwythos-9B is a full-parameter reasoning model post-trained on over 500 million tokens of high-quality Claude Mythos / Claude Fable traces with chain-of-thought generated in-house by Empero AI's internal `rethink` tool. It dominates the base Qwen3.5-9B under matched evaluation (**+34 pts MMLU, +30 pts gsm8k-strict, +19 pts gsm8k-flex**), supports **native function calling** per the Qwen3.5 spec, and ships with a **1,048,576-token (1M) context window** via YaRN rope-scaling enabled by default.

For full training details, evaluation numbers, and capability writeup, see the [**base model card**](#).

Files

Normal text weights — fixed v3 replacements

File	Quant	Size	Notes
<code>Qwythos-9B-Claude-Mythos-5-1M-Q4_K_M.gguf</code>	Q4_K_M	5.24 GiB / 5.63 GB	recommended default — fixed v3, best compatibility
<code>Qwythos-9B-Claude-Mythos-5-1M-Q5_K_M.gguf</code>	Q5_K_M	6.02 GiB / 6.47 GB	fixed v3, balanced quality / size
<code>Qwythos-9B-Claude-Mythos-5-1M-Q6_K.gguf</code>	Q6_K	6.85 GiB / 7.36 GB	fixed v3, high quality
<code>Qwythos-9B-Claude-Mythos-5-1M-Q8_0.gguf</code>	Q8_0	8.87 GiB / 9.53 GB	fixed v3, near-lossless
<code>Qwythos-9B-Claude-Mythos-5-1M-BF16.gguf</code>	BF16	16.69 GiB / 17.92 GB	fixed v3, full precision conversion base

If you don't know which to pick, **Q4_K_M is the right starting point** — it's the smallest practical quant with good quality preservation.

MTP-enabled text weights — fixed v3 variants

These include the restored Qwen3.5-compatible MTP head inside the GGUF. Use them with llama.cpp builds that support MTP draft speculation, for example `--spec-type draft-mtp`.

File	Quant	Size	Notes
<code>Qwythos-9B-Claude-Mythos-5-1M-MTP-Q4_K_M.gguf</code>	Q4_K_M + MTP	5.48 GiB / 5.89 GB	recommended MTP default
<code>Qwythos-9B-Claude-Mythos-5-1M-MTP-Q5_K_M.gguf</code>	Q5_K_M + MTP	6.26 GiB / 6.73 GB	MTP, balanced quality / size
<code>Qwythos-9B-Claude-Mythos-5-1M-MTP-Q6_K.gguf</code>	Q6_K + MTP	7.09 GiB / 7.62 GB	MTP, high quality
<code>Qwythos-9B-Claude-Mythos-5-1M-MTP-Q8_0.gguf</code>	Q8_0 + MTP	9.11 GiB / 9.79 GB	MTP, near-lossless
<code>Qwythos-9B-Claude-Mythos-5-1M-MTP-BF16.gguf</code>	BF16 + MTP	17.14 GiB / 18.41 GB	MTP, full precision conversion base

Vision projector — for image input

File	Size	Notes
<code>mmproj-Qwythos-9B-Claude-Mythos-5-1M-F16.gguf</code>	0.86 GiB / 0.92 GB	CLIP-style vision encoder + projector; required for images , pairs with any normal or MTP quant above

Qwythos inherits its **vision tower from the Qwen3.5-9B base model** — the vision path was *frozen* during SFT (training was text-only), so the vision behavior is identical to base Qwen3.5-9B's multimodal capability. The mmproj is interchangeable with any community-built Qwen3.5-9B `mmproj-*.gguf`.

Quick start

llama.cpp (`llama-cli`)

```
llama-cli \
  -m Qwythos-9B-Claude-Mythos-5-1M-Q4_K_M.gguf \
  -p "Walk through the biochemistry of how organophosphate nerve agents inhibit acetylcholinesterase." \
  -n 8192 \
  --temp 0.6 --top-p 0.95 --top-k 20 --repeat-penalty 1.05 \
  -c 16384
```

Ollama

```
ollama run hf.co/empero-ai/Qwythos-9B-Claude-Mythos-5-1M-GGUF:Q4_K_M
```

LM Studio / jan / KoboldCpp

Drop any of the `.gguf` files into your runtime's model directory. Qwythos uses the standard Qwen3.5 chat template; modern GGUF runtimes load it automatically from the file.

llama.cpp with MTP draft speculation

```
llama-server \
  -m Qwythos-9B-Claude-Mythos-5-1M-MTP-Q4_K_M.gguf \
  --spec-type draft-mtp \
  --spec-draft-n-max 6 \
  -c 16384 --port 8080
```

MTP support requires a recent llama.cpp build. If your runtime does not support MTP yet, use the normal fixed v3 files above.

Vision (image input)

Qwythos supports **image input** out of the box. Download both a text quant and the `mmproj-*.gguf` file from this repo, then run with llama.cpp's multimodal CLI or server.

llama.cpp (`llama-mtmd-cli`)

```
llama-mtmd-cli \
  -m Qwythos-9B-Claude-Mythos-5-1M-Q4_K_M.gguf \
  --mmproj mmproj-Qwythos-9B-Claude-Mythos-5-1M-F16.gguf \
  --image ./photo.jpg \
  -p "Describe this image in detail." \
```

```
--temp 0.6 --top-p 0.95 --top-k 20 \  
-c 16384
```

llama.cpp server (OpenAI-compatible API with images)

```
llama-server \  
-m Qwythos-9B-Claude-Mythos-5-1M-Q4_K_M.gguf \  
--mmproj mmproj-Qwythos-9B-Claude-Mythos-5-1M-F16.gguf \  
-c 16384 --port 8080
```

Then POST to `/v1/chat/completions` with an image URL or base64 payload — the standard OpenAI vision API shape works.

LM Studio

Load the text quant; LM Studio detects the matching `mmproj-*.gguf` in the same folder and enables the image-attach button automatically.

What vision unlocks

Since Qwythos inherits its vision tower unchanged from Qwen3.5-9B base, expect Qwen3.5-9B's documented vision capabilities: detailed image description, OCR (printed + handwritten), chart/table reading, UI/document understanding, basic spatial reasoning.

Honest note: the SFT used to produce Qwythos was **text-only** — we did not fine-tune the vision tower or train on any image-paired data. Image-grounded reasoning therefore inherits the base model's behavior; it has not been independently evaluated as part of this release. If your application is *primarily* vision-driven, validate on your own use case first.

Sampling recommendations

Qwythos is a reasoning model — every response opens with a `<think>...` block before the final answer. Use these settings as defaults:

Parameter	Value
<code>temperature</code>	0.6
<code>top_p</code>	0.95
<code>top_k</code>	20
<code>repeat_penalty</code>	1.05
<code>max_new_tokens</code>	16384 (generous budget for <code><think></code> + answer)

These match Qwen3.5's official thinking-mode recommendations. **Avoid greedy decoding and very-low-temperature sampling ($T \leq 0.3$)** — both can cause repetition loops on long reasoning generations.

Long context (1M tokens)

The GGUFs ship with YaRN rope-scaling baked in for a **1,048,576-token context window** (4× extension over the 262k native).

To use the full 1M window in `llama-cli`, set `-c 1010000` (or any context length up to that). For shorter prompts, lower `-c` to reduce KV-cache memory — at default settings llama.cpp will autosize.

A single H100/H200-class GPU comfortably handles **256k–512k**; the full 1M typically needs tensor-parallel multi-GPU or aggressive KV-cache offload.

Capabilities (from the base model card)

- **+34 pts MMLU, +30 pts gsm8k-strict, +19 pts gsm8k-flex** vs. base Qwen3.5-9B under matched lm-eval-harness evaluation
- **Native function calling** per Qwen3.5's chat-template spec — emits `<tool_call>` `<function=NAME><parameter=NAME>VAL</parameter></function></tool_call>` blocks ready for any tool-use loop
- **Self-correcting with tools:** in a 7-prompt tool-use harness (Python executor + DuckDuckGo search), Qwythos produced source-cited correct answers on 7/7, including 4/4 closed-book failure-modes from the original review
- **Uncensored** — engages seriously with technically demanding questions across cybersecurity, red-teaming, biology, pharmacology, and clinical

medicine

- **1,048,576-token (1M) context** — YaRN rope-scaling enabled by default

For full eval transcripts and per-task numbers, see the [base model card's evals/ folder](#).

Limitations

- **Reasoning model.** Every answer opens with a `<think>` block; allow generous `max_new_tokens` and parse/strip `<think>...</think>` for end users.
 - **Use recommended sampling.** Greedy / very-low-temp can cause repetition loops.
 - **Verify specifics in safety-critical contexts.** Like all closed-book LLMs in this weight class, Qwythos can over-commit to specific identifiers (CVEs, hashcat modes, drug positions) it isn't certain about. Pair with retrieval or function calling in such deployments — the model uses tools cleanly when offered them.
 - **Uncensored — add your own application-level review/safety layer** for end-user-facing deployments where that matters.
-

Acknowledgements

- Developed and released by Matuxyz in collab
- Base model: [Qwen3.5-9B](#) (Alibaba Qwen team)
- Quantization: [llama.cpp](#) (ggml-org)
- Vision projector (`mmpoj`): inherited from Qwen3.5-9B (vision tower unchanged); F16 GGUF re-hosted with thanks to [Unsloth](#) for the original conversion

Matuxy/Qwythos-9B-Claude-Mythos-5-1M-GGUF-bucket

